

Optimization

$$\text{CPU time} = \text{IC} \times \text{CPI} \times \text{CC}$$

- A chained if-then-else versus switch/case (use of Jumptable)
- Impact of procedure calls
 - Tail recursion
 - Inline
 - Hardware vs software
- Pipelining increases throughput (and equivalently reduces CPI or clock cycle)
 - Software resolution: Rescheduling
 - Register allocation
- Loops
 - unrolling: reduce loop overhead, enable rescheduling
 - Reduce inefficiency
 - re-ordering (take advantage of locality)

Handling multiple cases

use a series of chained conditional branches: if-then-else

if x = 1 then ...
else if x = 2 then ...
else if x = 3 then ...
....

Convert to assembly code

```
    @ assume x is stored in r0
    mov    r1, #1          @ r1 = 1
    cmp    r0, r1          @ r0
    bne    Else1

    ...                    @ code for when x = 1
Else1: mov    r1, #2
    cmp    r0, r1
    bne    Else2

    ...                    @ code for when x = 2
Else2: mov    r1, #3
    cmp    r0, r1
    bne    Else3

    ...
```

Drawbacks: Given x, its value has to be compared multiple times until a match is found, this can be costly when the chain is long and the match happens to be located towards the end of the chain.

use switch/case

```
int switchexample (int x) {  
    int y;  
    switch(x) {  
        case 1: y = x; break;  
        case 2: y = 2 * x;          /* fall through */  
        case 3: y = 3 * x; break;  
                                   /* missing case 4 */  
        case 5: y = 5 * x; break;  
        default: y = x; break;  
    }  
    return y;  
}
```

Implementations in assembly using a jumptable (see next slide)

Recall that this jumptable technique is also used in object oriented programming for handling methods for objects (Vtable).

Switchexample:

```
    ldr    r2, =Jtable      @ address of Jtable is in r2
    mov    r1, #0
    cmp    r0, r1           @ check if x (which is in r0) is smaller than 0
    blt    Default
    mov    r1, #5
    cmp    r0, r1           @ check if x is greater than 5
    bgt    Default
    add    r2, r2, r0, LSL #2 @ x multiply by 4 is used as the offset to r2
                                @ r2 points to the correct element in JTable
```

Default:

```
    ldr    r2, [r2]         @ load corresponding Jtable element, which
                                @ itself is an address of the case label.
```

```
    mov    pc, r2
```

```
Case1:  mov    r7, r0
        b      done
```

```
Case2:  mov    r7, r0, lsl #1 @ fall through to case 3
```

```
Case3:  mov    r3, #3
        mul    r7, r3, r0
        b      done
```

```
Case4:  b      Case0         @ undefined case is treated as default
```

```
Case5:  mov    r3, #5
        mul    r7, r3, r0
        b      done
```

```
Case0:  mov    r7, r0         @default case
        b      done
```

```
done:   mov    pc, lr
```

.data

```
JTable: .word Case0, Case1, Case2, Case3, Case4, Case5
```

Drawbacks: When the range is large and cases are sparse, the jumtable becomes not space economical. A chained if-then-else is a better option for such situations.

e.g.,

```
int switchexample (int x) {  
    int y;  
    switch(x) {  
        case 1: y = x; break;  
        case 20: y = 20 * x;          /* fall through */  
        case 330: y = 330 * x; break;  
        case 5000: y = 5000 * x; break;  
        default: y = x; break;  
    }  
    return y;  
}
```

Note: the switch/case is introduced as a distinct construct in the C language so that a compiler can “know” a jumtable implementation is preferred to a chained if-then-else implementation.

Procedure calls

Optimization to reduce the overhead (maintaining the stack)

-Inlining (cut & paste)

Examples:

```
inline int max (int a, int b)
{
    if (a > b)
        return a;
    else
        return b;
}

a = max (x, y);

/*
   This is now equivalent to

   if (x > y)
       a = x;
   else
       a = y;
*/
```

pros: reduce run-time cost;

allow for more transformations conducive to optimization.

cons: increase code size, thus less desirable in embedded systems where memory is limited.

Procedure calls

Optimization to reduce the overhead (maintaining the stack)

-Tail recursion:

- when no further computation follows a recursive call;
- can be easily converted to iteration (automatically by the compiler);
- no stack overhead is needed any more.


```

gcd(a, b) {
  while (a!=b) {
    if(a>b) a = a - b;
    else b = b - a;
  }
  return a;
}

```

gcd:

```

Loop:    cmp    r0, r1
         beq    Return

If:      ble    Else
         sub    r0, r0, r1
         b      Loop
Else:    sub    r1, r1, r0
         j      Loop

Return:  mov    r7, r0
         mov    pc, lr

```

```

gcd(a, b) {
  if(a==b) return a;
  else if (a>b) return gcd(a-b, b);
  else return gcd(a, b-a);
}

```

gcd:

sub	sp, sp, 4	@ push
str	lr, [sp, #0]	@push

```

cmp    r0, r1
beq    Return

If:      ble    Else
         sub    r0, r0, r1
         b      Rec
Else:    sub    r1, r1, r0
Rec:     bl     gcd      @ recursive call

```

@ Note that here is the return point after the recursive call,
 @ but there is no more instructions to execute here other than fall
 @ though to the end of the whole procedure !
 @ Therefore, no need to save this return addr onto the stack.
 @ So, remove the code for push and pop, and change bl to b.

ldr	lr, [sp, #0]	@ pop
add	sp, sp, #4	@ pop

```

move    r7, r0
mov     pc, lr

```

fact(n)

if n = 1 return 1

else return n * fact(n-1)

“tailize” →

fact1(n, p)

if n = 1 return p

else return fact1(n-1, n*p)

@assume n is in r0

fact:

```
sub    sp, sp, #8
str    lr, [sp, #4]
str    r0, [sp, #0]
mov    r2, #1
cmp    r0, r2
bne    Else
mov    r7, #1
b      Return
```

> push

Else: sub r0, r0, 1 @ n = (n-1)

```
bl     fact
ldr    r0, [sp, #0]
mul    r7, r0, r7
```

Return:

```
ldr    lr, [sp, #4]
add    sp, sp, #8
mov    pc, lr
```

~ pop

@ assume n is in r0, and p is in r1

fact1:

```
mov    r2, #1
cmp    r0, r2
bne    Else
mov    r7, r1
b      Return
```

~~~~~

Else: mul r1, r1, r0

```
sub    r0, r0, 1
b      fact1
```

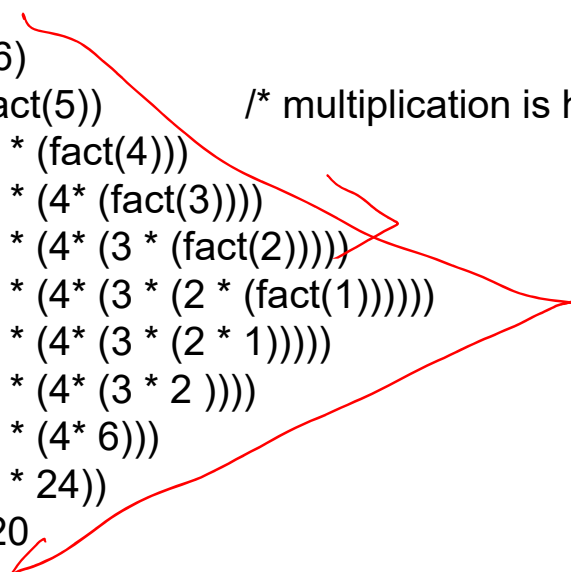
Return:

```
mov    pc, lr
```

~~~~~

The difference can be observed via the semantics at the high level programming language


```
fact(6)
= 6 * (fact(5))      /* multiplication is held off because fact(5) is not known */
= 6 * (5 * (fact(4)))
= 6 * (5 * (4 * (fact(3))))
= 6 * (5 * (4 * (3 * (fact(2)))))
= 6 * (5 * (4 * (3 * (2 * (fact(1))))))
= 6 * (5 * (4 * (3 * (2 * 1))))
= 6 * (5 * (4 * (3 * 2)))
= 6 * (5 * (4 * 6))
= 6 * (5 * 24)
= 6 * 120
= 720
```



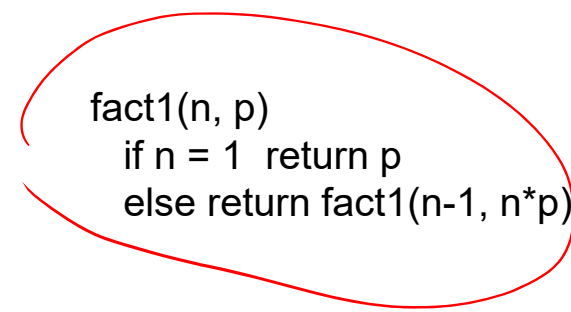
→ Correspond to the growth and shrink of the stack in memory

n p

```
fact1(6, 1)
= fact1(5, 6)
= fact1(4, 30)
= fact1(3, 120)
= fact1(2, 360)
= fact1(1, 720)
```



```
fact1(n, p)
  if n = 1 return p
  else return fact1(n-1, n*p)
```



Reduce loop inefficiency

Loop invariant: instructions whose result does not change from iteration to iteration, and therefore can be moved outside the loop without affecting the semantics of the program.

For example,

```
for(i = 0; i < n; i++) {  
    z = x + y;  
    a[i] = 4*i + z*z;  
}
```

Loop-invariant: $z = x + y$, and $z*z$

```
z = x + y;  
w = z * z;  
for(i = 0; i < n; i++) {  
    a[i] = 4*i + w;  
}
```

Strength reduction: a costly operation is replaced with an equivalent by less expensive operation

For example,

Multiplication and division by a power of 2 can be achieved with shifting

$4*i$ by $i \ll 2$

Reduce loop inefficiency

```
1 /* Convert string to lower case: slow */
2 void lower1(char *s)
3 {
4     int i;
5
6     for (i = 0; i < strlen(s); i++)
7         if (s[i] >= 'A' && s[i] <= 'Z')
8             s[i] -= ('A' - 'a');
9 }
10
```

Credit: Bryant & O'Hallaron, CSAPP

Credit: Bryant & O'Hallaron, CSAPP

Unrolling loops & rescheduling

Example: function that computes the dot product of two n-element int arrays.

```
dp = 0;
for(i=1; i<n; i++) {
    dp += A[i] x B[i];
}
```

```
1      mov      r3, r0   @ r0 = &A
2      mov      r4, r1   @ r1 = &B
3      mov      r5, r2   @ r2 = n
3      mov      r9, #0
5      b L2
L1:    ldr      r6,[ r3, #0]
7      ldr      r7, [r4, #0]
8      mul      r8, r6, r7
9      add      r9, r9, r8
10     add      r3, r3, #4
11     add      r4, r4, #4
12     sub      r5, r5, #1
13     mov      r10, #0
14     cmp      r5, r10
L2:    bne      L1
```

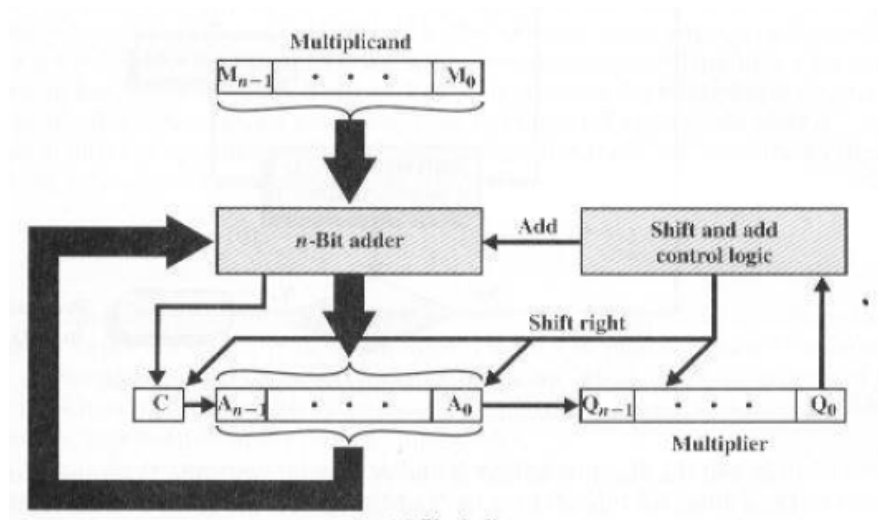
Let's assume the pipelined CPU has one cycle delay for load-and-use dependency, and 3 cycle delay for read-before-write for mul. Without rescheduling, this program needs to have delayed 4 cycles per iteration. E.g., if $n=200$, idled cycles = $4 \times 200 = 800$.

If we reschedule as follows, we can remove one delayed cycle after ldr r7, two delayed cycles after mul, and leave only one delayed cycle after mul.

original			After rescheduling		
<hr/>			<hr/>		
1	mov	r3, r0 @ r0 = &A	1	mov	r3, r0 @ r0 = r&A
2	mov	r4, r1 @ r1 = &B	2	mov	r4, r1 @ r1 = &B
3	mov	r5, r2 @ r2 = n	3	mov	r5, r2 @ r2 = n
3	mov	r9, #0	3	mov	r9, #0
5	b	L2	5	b	L2
L1:	ldr	r6, [r3, #0]	L1:	ldr	r6, [r3, #0]
7	ldr	r7, [r4, #0]	7	ldr	r7, [r4, #0]
8	mul	r8, r6, r7	8	sub	r5, r5, #1
9	add	r9, r9, r8	9	mul	r9, r6, r7
10	add	r3, r3, #4	10	add	r3, r3, #4
11	add	r4, r4, #4	11	add	r4, r4, #4
12	sub	r5, r5, #1	12	add	r9, r9, r8
13	mov	r10, #0	13	mov	r10, #0
14	cmp	r5, r10	14	cmp	r5, r10
L2:	bne	L1	L2:	bne	L1

Hardware (e.g., multiplication)
mul r1, r2, r3

The function can natively supported by an instruction at the machine level with hardware implementation as shown below. While this hardware solution seems to be the best -- neither overhead run-time cost like for subroutine nor increased code size of inline function, this leads to more expensive hardware and therefore should only be resorted to for optimizing mostly common functions such as multiplication.



Another example of hardware solution to optimization.

The instruction

```
...  
bl      sub1  
...
```

can be simulated with a couple ARM instructions as follows.

```
...  
ldr     lr, =Return_here  
b       sub1
```

Return_here: ...

Design philosophy of RISC v.s. CISC:

CISC favors hardware solution, i.e, more instructions natively supported by hardware, whereas RISC favors fewer instructions and provides pseudoinstructions supported by the assembler to ease the programming.

Reordering nested loops to increase cache hits (via locality)

```
for i := 1 to n
  for j := 1 to n
    A[i, j] := 0
```

If A is laid out in row-major order, and if each cache line contains m elements of A, then this code will suffer n^2/m cache misses. On the other hand, if A is laid out in column-major order, and if the cache is too small to hold n lines of A, then the code will suffer n^2 misses, fetching the entire array from memory m times. The difference can have an enormous impact on performance. A loop-reordering compiler can improve this code by *interchanging* the nested loops:

```
for j := 1 to n
  for i := 1 to n
    A[i, j] := 0
```

In more complicated examples, interchanging loops may improve locality of reference in one array, but worsen it in others. Consider this code to transpose a two-dimensional matrix:

```
for j := 1 to n
  for i := 1 to n
    A[i, j] := B[j, i]
```